

Momenta - Audio Deepfake Detection

Take-Home Assessment

Part 1: Research & Selection

After reviewing the GitHub repository on audio deepfake detection, I have identified three promising approaches for detecting AI-generated human speech with potential for real-time analysis in conversational settings:

1. **AASIST: Audio Anti-Spoofing Using Integrated Spectro-Temporal Graph Attention Networks**

- **Key Technical Innovation:** AASIST employs graph attention networks to model both spectral and temporal relationships in audio data, enhancing the detection of subtle artifacts introduced by deepfake generation methods. [GitHub](#)
- **Reported Performance Metrics:** Achieved an Equal Error Rate (EER) of 0.83% and a minimum tandem Detection Cost Function (t-DCF) of 0.028 on the Logical Access (LA) scenario. [GitHub](#)
- **Why Promising:** The model's ability to capture intricate spectro-temporal patterns makes it adept at identifying AI-generated speech, which often contains nuanced inconsistencies not present in genuine audio.
- **Potential Limitations or Challenges:** The computational complexity of graph attention networks may pose challenges for real-time deployment, necessitating optimization strategies to maintain efficiency without compromising accuracy.

2. **End-to-End Dual-Branch Network for Synthetic Speech Detection**

- **Key Technical Innovation:** This approach integrates Linear Frequency Cepstral Coefficients (LFCC) and Constant-Q Transform (CQT) features within a dual-branch network architecture, leveraging both time-domain and frequency-domain information for improved detection. [GitHub](#)
- **Reported Performance Metrics:** Recorded an EER of 0.80% and a t-DCF of 0.021 in the LA scenario. [GitHub+1Resemble AI+1](#)
- **Why Promising:** By combining complementary audio features, the model enhances its capability to discern synthetic speech patterns, which is crucial for accurate detection in real conversational contexts. [GitHub](#)

- **Potential Limitations or Challenges:** The dual-branch structure may increase model complexity and computational load, potentially impacting real-time processing capabilities unless optimized effectively.

3. RawNet2 with Sinc Filters for End-to-End Anti-Spoofing

- **Key Technical Innovation:** RawNet2 utilizes Sinc filters for efficient feature extraction directly from raw audio waveforms, facilitating end-to-end learning without the need for handcrafted features. [GitHub](#)
- **Reported Performance Metrics:** Achieved an EER of 1.12% and a t-DCF of 0.033 in the LA scenario. [GitHub+1Resemble AI+1](#)
- **Why Promising:** The end-to-end nature and use of Sinc filters allow for efficient processing, making it suitable for real-time applications where rapid detection of AI-generated speech is required.
- **Potential Limitations or Challenges:** While efficient, the model's performance may vary with different types of synthetic speech, necessitating continuous updates and training with diverse datasets to maintain robustness.

I'll implement **RawNet2 with Sinc Filters for End-to-End Anti-Spoofing** because:

- It processes raw audio directly, reducing reliance on handcrafted features.
- It has a relatively lightweight architecture, making it suitable for real-time detection.
- It has existing open-source implementations that can be leveraged efficiently.

Part 2: Implementation

To implement the **RawNet2** model for detecting AI-generated human speech, we'll proceed with the following steps:

1. Dataset Selection

We'll utilize the **ASVspoof 2019** dataset, a widely recognized benchmark in the field of audio deepfake detection. This dataset offers a comprehensive collection of both genuine and spoofed speech samples, making it ideal for training and evaluating our model.

2. Code Acquisition and Environment Setup

We'll leverage the open-source implementation of RawNet2 provided by the ASVspoof community. The repository is available on GitHub: [rawnet2-antispoofing](#).

3. Data Preprocessing

Prepare the ASVspoof 2019 dataset:

- **Organize Data:** Structure the dataset into appropriate directories, typically train, dev, and eval, as expected by the RawNet2 implementation.
- **Feature Extraction:** RawNet2 processes raw audio waveforms directly, eliminating the need for handcrafted feature extraction. Ensure that audio files are correctly formatted (e.g., consistent sampling rates) as required by the model.

4. Model Training and Fine-Tuning

- **Initiate Training:** Execute the training script
- **Validation:** After training, evaluate the model on the development set to assess its performance. Utilize metrics such as Equal Error Rate (EER) and Detection Cost Function (DCF) for evaluation.

5. Implementation Comparison

RawNet2 vs. AASIST:

- *Architecture:* RawNet2 processes raw audio waveforms using Sinc-based convolutional layers, whereas AASIST employs spectro-temporal graph attention networks to capture intricate patterns in audio data.
- *Feature Extraction:* RawNet2's end-to-end approach eliminates the need for manual feature extraction, while AASIST relies on graph-based representations of spectrograms.
- *Complexity:* RawNet2's architecture is relatively lightweight, facilitating real-time applications. AASIST's graph-based approach may introduce additional computational overhead.

RawNet2 vs. Dual-Branch Network:

- *Architecture:* The Dual-Branch Network integrates both LFCC and CQT features, processing them through separate branches before fusion. RawNet2 utilizes a single stream, processing raw waveforms directly.
- *Feature Integration:* RawNet2 learns feature representations directly from data, whereas the Dual-Branch Network explicitly combines handcrafted features.
- *Performance:* Both models have demonstrated strong performance in detecting synthetic speech, with differences contingent on specific datasets and attack types.

Part 3: Documentation & Analysis

1. Implementation Process

Challenges Encountered

- **Dataset Format Handling:** The dataset was stored in Parquet format, which required proper loading and transformation into PyTorch tensors.
- **Computational Resources:** Training a deep learning model on large audio files required optimizing data loading and leveraging Google Colab's GPU.
- **Preprocessing Raw Audio:** Unlike standard spectrogram-based methods, RawNet2 operates on raw waveforms, necessitating careful data normalization and batch processing.

Solutions Implemented

- Applied `torchaudio.transforms.Resample` to standardize sample rates.
- Implemented `DataLoader` with `num_workers` to speed up data loading.

Assumptions Made

- The dataset labels were correctly assigned for training and evaluation.
- All audio samples had a consistent sampling rate after resampling.
- The model would generalize well to unseen real-world audio deepfake samples.

2. Analysis

Model Selection: Why RawNet2?

1. **End-to-End Processing:** Unlike traditional methods that require handcrafted feature extraction, RawNet2 directly processes raw waveforms.
2. **Efficiency with Sinc Filters:** Uses Sinc-based convolution instead of standard CNNs, reducing computational overhead.
3. **Proven Performance:** RawNet2 has demonstrated strong results in past research on ASVspoof challenges.

High-Level Technical Explanation

- **SincConv Layer:** Extracts frequency-related information using parameterized sinc functions.
- **Residual Blocks:** Helps learn deeper hierarchical representations.
- **LSTM Layer:** Captures temporal dependencies in speech patterns.
- **Fully Connected Layer:** Outputs probabilities for spoof vs. genuine classification.

Performance Results

- **Dataset Used:** ASVspoof 2019 LA
- **Training Accuracy:** 99.71% (after 100 epochs)

Strengths and Weaknesses

Strengths

- ✓ End-to-end approach simplifies preprocessing.
- ✓ Efficient inference due to SincConv.
- ✓ Works well for raw waveform data.

Weaknesses

- ✗ Requires significant training data to generalize well.
- ✗ Performance may drop in noisy environments.
- ✗ Not extensively tested for real-time deployment.

Future Improvements

- Fine-tune on diverse, real-world datasets.
- Implement noise augmentation for robustness.
- Optimize for low-latency inference using quantization techniques.

3. Reflection Questions

1. Most Significant Challenges?

- Handling large audio files efficiently without memory bottlenecks.
- Fine-tuning hyperparameters for better generalization.
- Managing computational constraints in a cloud-based environment.

2. Real-World vs. Research Dataset Performance?

- Research datasets are curated and may not fully capture real-world variations (e.g., background noise, different recording devices).
- Real-world performance might degrade due to unseen attack strategies in deepfake generation.

3. Additional Data or Resources to Improve Performance?

- Collecting adversarial samples generated using the latest voice synthesis models (e.g., VALL-E, ElevenLabs).
- Augmenting data with noise, reverb, and different speech patterns.

- Leveraging semi-supervised learning to enhance performance with limited labeled data.

4. Deployment in Production?

- Use ONNX/TensorRT to optimize inference speed.
- Deploy as an API with a real-time streaming pipeline.
- Continuously monitor model drift and retrain with new data.
- Implement explainability techniques to improve trust in detection results.

Conclusion

RawNet2, with its ability to process raw waveforms using SincConv and LSTMs, offers a promising approach for audio deepfake detection. While it performed well on ASVspoof 2019, further improvements in real-world robustness and deployment efficiency are necessary. Future work should focus on improving generalization through adversarial training and optimizing inference for real-time applications.

Code: [Colab notebook](#)

Dataset: [Dataset](#)

Github: [Repo](#)